

In [4]:

```
%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import random
```

4a

In [213]:

```
delta = 3
huber = lambda e, delta: ((e**2)/2) if np.abs(e) <= delta else ((delta * np.abs(
e)) - ((delta**2)/2))

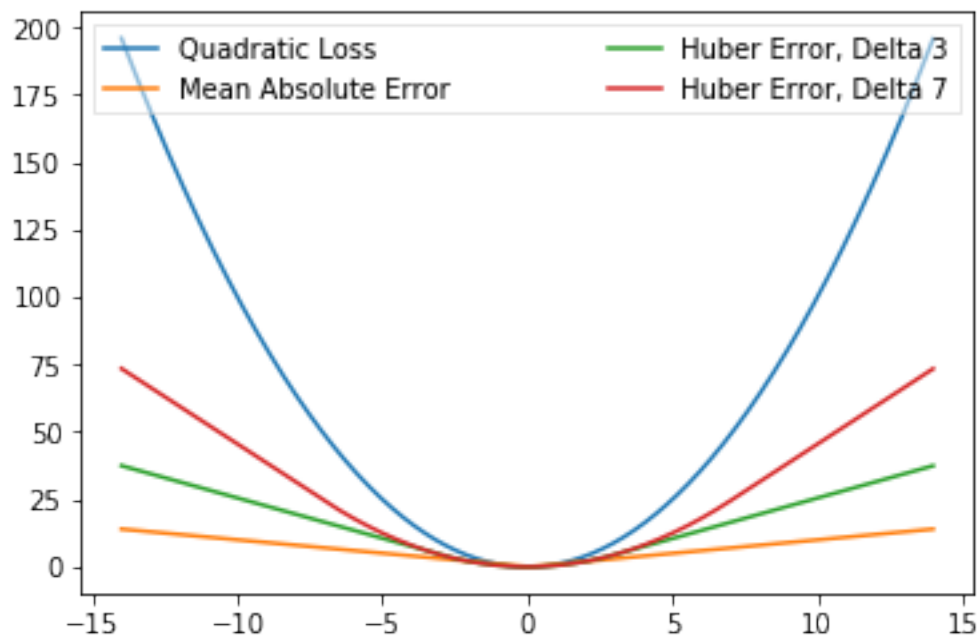
ax = plt.subplot(111)
t1 = np.arange(-14.0, 14.0, 0.01)

t2 = [huber(x, delta) for x in t1]

plt.plot(t1, t1**2, label="Quadratic Loss")
plt.plot(t1, np.abs(t1), label="Mean Absolute Error")
plt.plot(t1, [huber(x, 3) for x in t1], label="Huber Error, Delta 3")
plt.plot(t1, [huber(x, 7) for x in t1], label="Huber Error, Delta 7")

leg = plt.legend(loc='best', ncol=2, mode="expand", shadow=False, fancybox=False
)
leg.get_frame().set_alpha(0.5)

plt.show()
# plot second time
```



Relative advantages and disadvantages of loss functions for linear regression.

Given the plots and their mathematical equations, it appears that Mean Absolute Error is less sensitive to outliers than Quadratic Loss. A square of a single large outlier might be enough to alter our weights in a significant way. Therefore in noisy datasets with large outliers, it appears that MAE is a better choice. However, since mean absolute error doesn't reduce the slope as it approaches optima, it might be more sensitive to large learning rates and might never reach an optima.

Huber Loss on the other hand is a good tradeoff between the above two loss functions. It is linear at large loss values allowing it to be robust to outliers, while employing a nice curvature with a reducing slope as it approaches optima. The catch with huber loss, however, is that we may not necessarily know the delta beforehand.

4b

In [2]:

```
def genData(numPoints, slope, intercept):
    x = np.zeros(shape=(numPoints, 2))
    y = np.zeros(shape=numPoints)

    x[:,0] = 1
    x[:,1] = np.random.uniform(-2, 2, size = numPoints)

    for i in range(0, numPoints): y[i] = (x[i][1] * slope) + (intercept + np.random.normal(0,4))

    return x, y
```

In [207]:

```
def gradientDescent(x, y, theta, alpha, m, numIterations, loss_type):

    xTrans = x.transpose()
    for i in range(0, numIterations):

        hypothesis = np.dot(x, theta)

        if loss_type == 'quadratic':

            loss = hypothesis - y
            cost = np.sum(loss ** 2) / (2 * m)

        elif loss_type == 'mean_absolute':

            loss = hypothesis - y

            loss[loss>0] = 1
            loss[loss<0] = -1

            cost = np.sum(loss) / (2 * m)

        elif loss_type == 'huber':

            delta = 2
            huber = lambda e, delta: ((e/2)**2) if np.abs(e) <= delta else ((delta * np.abs(e)) - ((delta**2)/2))
            loss = hypothesis - y
            loss[abs(loss) <= delta] = loss[abs(loss) <= delta]
            loss[abs(loss) > delta] = delta * np.sign(loss[abs(loss) > delta])
            cost = huber(np.sum(loss), delta) / (2 * m)

        #print("\n\nIteration %d / Cost: %f" % (i, cost))

        gradient = np.dot(xTrans, loss) / m #Partial derivative
        theta = theta - alpha * gradient

        #print ('gradient' + str(gradient))
        #print ('theta' + str(theta))

    return theta
```

In [221]:

```
#Referenced Stackoverflow for this
```

```
from numpy import array, dot, transpose
from numpy.linalg import inv

def analytical_linear_regression(x_train, y_train):

    X = np.array(x_train)
    y = np.array(y_train)

    Xt = transpose(X)
    product = dot(Xt, X)
    theInverse = inv(product)
    w = dot(dot(theInverse, Xt), y)

    return w
```

In [222]:

```
analytical_linear_regression(x,y)
```

Out[222]:

```
array([3.23055664, 4.20992994])
```

In [203]:

```
def gradientDescentSto(x, y, theta, alpha, m, numIterations, loss_type, delta=2,
conv=0.00000001):

    xTrans = x.transpose()
    count = 0
    theta_prev = theta + delta + 1

    while (count < numIterations):

        count += 1

        for i in range(m):

            hypothesis = x[i][0] * theta[0] + x[i][1]*theta[1]

            if loss_type == 'quadratic': loss = hypothesis - y[i]

            if loss_type == 'mean_absolute':

                loss = hypothesis - y[i]
                loss = 1 if loss >= 0 else -1

            if loss_type == 'huber':

                loss = hypothesis - y[i]
                if loss < -delta: loss = -delta
                elif -delta <= loss and loss <= delta: loss = loss
                elif delta < loss: loss = delta

            gradient1 = x[i][0] * loss
            gradient0 = x[i][1] * loss

            theta_prev = theta
            theta[1] = theta[1] - alpha * gradient0
            theta[0] = theta[0] - alpha * gradient1


    return theta
```

Calling and Plotting

In [206]:

```
x, y = genData(100, 3, 2)
m, n = np.shape(x)

theta = gradientDescentSto(x, y, np.ones(n), alpha = 0.0001, m, numIterations =
5000, 'huber', 7)
hypothesis = lambda x: x * theta[1] + theta[0]
plt.plot([i for i in range(-10,10)], [hypothesis(i) for i in range(-10,10)])
plt.scatter([i[1] for i in x], y, s=10)
plt.show()
```



Call 1000 times and histogram for Huber, MAE and Analytical

In [492]:

```
thetas_huber = []
thetas_mean_abs = []
thetas_analytical = []

for i in range(100):
    #print (i)
    x, y = genData(100, 3, 2)
    m, n = np.shape(x)
    numIterations= 5000


    alpha = 0.0001
    thetas_huber.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'huber', 7))
    thetas_mean_abs.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'mean_absolute'))
    thetas_analytical.append(analytical_linear_regression(x,y))
```

In [411]:

```
plt.hist(np.array(thetas_huber)[: ,1], label = 'huber')
plt.hist(np.array(thetas_mean_abs)[: ,1], label = 'mean', alpha = 0.5)
plt.hist(np.array(thetas_analytical)[: ,1], label = 'analy', alpha = 0.5)

plt.plot([3 for i in range (25)], [i for i in range(25)])
#plt.plot(x,3,linestyle='solid')

plt.legend()
plt.show()
```




In [412]:

```
plt.hist(np.array(thetas_huber)[: ,0], label = 'huber')
plt.hist(np.array(thetas_mean_abs)[: ,0], label = 'mean', alpha = 0.5)
plt.hist(np.array(thetas_analytical)[: ,0], label = 'analy', alpha = 0.5)

plt.plot([2 for i in range (25)], [i for i in range(25)])

plt.legend()
plt.show()
```



Choice of Loss Function Affecting Slope and Intercept

As explained in 4a, Hubers loss function appears to have the highest frequency of slope closest to 3 and intercept with analytical coming in a close second. However, since analytical solution is the actual solution to our simulated values, that should be used as benchmark. As such, Huber outperforms MAE in estimating the slope and intercept.

Call 1000 times and histogram for Batch, Stochastic and Analytical

In [228]:

```
thetas_sto_quadratic = []
thetas_batch_quadratic = []
thetas_analytical = []

for i in range(100):
    #print (i)
    x, y = genData(100, 3, 2)
    m, n = np.shape(x)
    numIterations= 5000
    alpha = 0.001


    thetas_sto_quadratic.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'quadratic'))
    thetas_batch_quadratic.append(gradientDescent(x, y, np.ones(n), alpha, m, numIterations, 'quadratic'))
    thetas_analytical.append(analytical_linear_regression(x,y))
```

In [233]:

```
plt.hist(np.array(thetas_quadratic)[: ,0], label = 'Stochastic', alpha = 1)
plt.hist(np.array(thetas_batch_quadratic)[: ,0], label = 'Batch', alpha = 0.6)
plt.hist(np.array(thetas_analytical)[: ,0], label = 'Analytical', alpha = 0.3)
```

```
plt.plot([2 for i in range (25)], [i for i in range(25)])
```

```
plt.legend()
plt.show()
```




In [242]:

```
plt.hist(np.array(thetas_quadratic)[: ,1], label = 'Stochastic', alpha = 0.5)
plt.hist(np.array(thetas_analytical)[: ,1], label = 'Analytical', alpha = 0.7)
plt.hist(np.array(thetas_batch_quadratic)[: ,1], label = 'Batch', alpha = 1)

plt.plot([3 for i in range (25)], [i for i in range(25)])

plt.legend()
plt.show()
```



Choice of Algorithm Affecting Slope and Intercept

Using the analytical solution as the benchmark, it appears like batch almost exactly mimics the analytical solutions to the simulated data, outperforming stochastic in this case.

Add Anomalies

In [486]:

```
outlier_indices = [int(i) for i in np.random.uniform(0,100, size = 10)]
outlier_indices_indices = [int(i) for i in np.random.uniform(0,10, size = 5)]
x[outlier_indices]
```

Out[486]:

```
array([[ 1.          ,  1.40664088],
       [ 1.          ,  1.81175379],
       [ 1.          , -0.26609268],
       [ 1.          ,  0.57753574],
       [ 1.          ,  0.41188323],
       [ 1.          ,  0.03637198],
       [ 1.          ,  1.19956399],
       [ 1.          , -0.83287292],
       [ 1.          ,  0.7396422 ],
       [ 1.          , -0.43608485]])
```

In [488]:

```
positive_indices = [outlier_indices[i] for i in outlier_indices_indices]
negative_indices = [outlier_indices[i] for i in [outlier_indices.index(i) for i
in outlier_indices if outlier_indices.index(i) not in outlier_indices_indices]]
```

In [489]:

```
for i in positive_indices: x[i] += 0.5*x[i]
for i in negative_indices: x[i] -= 0.5*x[i]
```

Rerun 1000 times with anomalies

In [243]:

```
thetas_quadratic = []
thetas_quadratic.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'quadratic'))
```

In [245]:

```
thetas_huber = []
thetas_mean_abs = []
thetas_quadratic = []
thetas_analytical = []

for i in range(100):
    #print (i)
    x, y = genData(100, 3, 2)
    m, n = np.shape(x)
    numIterations= 5000
    alpha = 0.0001


    thetas_huber.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'huber', 7))
    thetas_mean_abs.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'mean_absolute'))
    thetas_quadratic.append(gradientDescentSto(x, y, np.ones(n), alpha, m, numIterations, 'quadratic'))
    thetas_analytical.append(analytical_linear_regression(x,y))
```

In [251]:

```
plt.hist(np.array(thetas_huber)[: ,1], label = 'huber')
plt.hist(np.array(thetas_mean_abs)[: ,1], label = 'mean', alpha = 0.5)
plt.hist(np.array(thetas_quadratic)[: ,1], label = 'quadratic', alpha = 0.5)

plt.plot([3 for i in range (25)], [i for i in range(25)])

#plt.plot(x,3,linestyle='solid')
plt.legend()
plt.show()
```




```
In [250]:
```

```
plt.hist(np.array(thetas_huber)[: ,0], label = 'huber')
plt.hist(np.array(thetas_mean_abs)[: ,0], label = 'mean', alpha = 0.5)
plt.hist(np.array(thetas_quadratic)[: ,0], label = 'quadratic', alpha = 0.5)

plt.plot([2 for i in range (25)], [i for i in range(25)])

plt.legend()
plt.show()
```



Choice of Loss in Noisy Data Affecting Slope and Intercept

Using the analytical solution as the benchmark, huber and quadratic outperform MAE in this case. This could be since the anomalies built into the data are balanced by design and therefore don't skew the squared error in any particular direction.

1)
i) The conditional probability of Y_1 and Y_2 is

$$f_{Y_1|Y_2}(Y_1|Y_2) = \frac{1}{\sqrt{2\pi(1-\rho^2)}\sigma_1} e^{\left(\frac{-(Y_1-a)^2}{2\sigma^2(1-\rho^2)}\right)}$$

$$\text{where } a = \mu_1 + \rho \frac{\sigma_1}{\sigma_2} (Y_2 - \mu_2).$$

~~The~~ conditional prob

The expected value of $Y_1|Y_2$ is

$$E[Y_1|Y_2] = \mu_1 + \rho \frac{\sigma_1}{\sigma_2} (Y_2 - \mu_2)$$

$$G[Y_1|Y_2] = \left[\mu_1 - \mu_2 \rho \frac{\sigma_1}{\sigma_2} \right] + \rho \frac{\sigma_1}{\sigma_2} Y_2$$

This is a linear combination of Y_2 , which mimics linear regression with one predictor. The variance is defined

$$\text{by } \sigma[Y_1|Y_2] = \sigma^2(1-\rho^2)$$

1

a) Joint probability distribution of Y_1 and Y_2 is

$$f_{Y_1, Y_2}(y_1, y_2) = \frac{1}{2\pi\sigma_1\sigma_2\sqrt{1-\rho^2}} \cdot e^{-\frac{1}{2(1-\rho^2)}$$

$$\left[\left(\frac{y_1 - \mu_1}{\sigma_1} \right)^2 + \left(\frac{y_2 - \mu_2}{\sigma_2} \right)^2 - 2\rho \left(\frac{y_1 - \mu_1}{\sigma_1} \right) \left(\frac{y_2 - \mu_2}{\sigma_2} \right) \right]$$

b) ~~margin~~

$$f_{Y_1|Y_2}(y_1, y_2) = \frac{f_{Y_1, Y_2}(y_1, y_2)}{f_{Y_2}(y_2)}$$

$f_{Y_2}(y_2)$ (marginal distribution) is.

$$\int_{-\infty}^{\infty} f_{Y_1, Y_2}(y_1, y_2) dy_1 = \frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(y_2 - \mu_2)^2}{2\sigma_2^2}}$$

Applying Bayes theorem we get.

$$f_{Y_1, Y_2}(Y_1, Y_2) = \frac{1}{\sqrt{2\pi}\sigma_1\sigma_2\sqrt{1-\rho^2}} \cdot e^{-\frac{1}{2(1-\rho^2)}} \left[\left(\frac{Y_1 - \mu_1}{\sigma_1} \right)^2 - 2\rho \left(\frac{Y_1 - \mu_1}{\sigma_1} \right) \left(\frac{Y_2 - \mu_2}{\sigma_2} \right) + \left(\frac{Y_2 - \mu_2}{\sigma_2} \right)^2 \right]$$

$$\frac{1}{\sqrt{2\pi}\sigma_2} e^{-\frac{(Y_2 - \mu_2)^2}{2\sigma_2^2}}$$

$$= \frac{1}{\sqrt{2\pi(1-\rho^2)}\sigma_1} e^{-\frac{(Y_1 - (\mu_1 + \frac{\rho\sigma_1}{\sigma_2}(Y_2 - \mu_2)))^2}{2\sigma_1^2(1-\rho^2)}}$$

The conditional probability of $Y_1 | Y_2$ is normally distributed having mean $\mu_1 + \rho \frac{\sigma_1}{\sigma_2} (Y_2 - \mu_2)$ and $\sigma^2 = (1-\rho^2)\sigma_1^2$

2
a) Quadratic loss is defined by.

$$L = \sum_{i=1}^N e_i^2.$$

Taking second derivative.

$$\frac{\partial L}{\partial e} = 2e ; \quad \frac{\partial^2 L}{\partial e^2} = 2$$

Since the $\frac{\partial^2 L}{\partial e^2} > 0$ this loss is convex.

b) L₁ norm $L = \sum_{i=1}^N |e_i|$

$$\frac{\partial L}{\partial e} = \text{sgn}(e) ; \quad \frac{\partial^2 L}{\partial e^2} = 0$$

Since $\frac{\partial^2 L}{\partial e^2} = 0$ this loss function is also convex.

$$c) \quad L = \sum_{i=1}^n l(e_i) \quad \text{where } l(e) = \begin{cases} \frac{1}{2}e^2 & \text{if } |e| \leq \delta \\ \delta|e| - \frac{1}{2}\delta^2 & \text{if } |e| > \delta \end{cases}$$

$$\frac{\partial L}{\partial e} = \begin{cases} e & \text{if } |e| \leq \delta \\ \delta \operatorname{sgn}(e) & \text{if } |e| > \delta \end{cases}$$

$$\frac{\partial^2 L}{\partial e^2} = \begin{cases} 1 & \text{if } |e| \leq \delta \\ 0 & \text{if } |e| > \delta \end{cases}, \quad \text{both of which are greater than or equal to 0.}$$

Therefore this function is convex as well.

3

$$a) \quad y_i = \theta_0 + \theta_1 x_i + e_i, \quad i = 1, \dots, N.$$

Our goal is to minimise the cost function

$$J(\theta_0, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2.$$

where m is the number of datapoints

Our hypothesis is ~~$h(x)$~~ $h(x) = \theta_0 x_0 + \theta_1 x_1 + \dots + \theta_n x_n$

$$h(x) = \theta^T x.$$

$$\text{where, } \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_n \end{bmatrix} \quad \text{and} \quad x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix}$$

The cost function converted into matrix notation, we get.

$$J(\theta) = \frac{1}{2m} (X\theta - y)^T (X\theta - y)$$

$$J(\theta) = \frac{(X\theta)^T - y^T (X\theta - y)}{2m}$$

$$= \frac{(X\theta)^T X\theta - (X\theta)^T y - y^T (X\theta) + y^T y}{2m}$$

$$= \frac{\theta^T X^T X\theta - 2(X\theta)^T y + y^T y}{2m}$$

Differentiating and setting it to 0

$$\frac{\partial}{\partial \theta} \left(\frac{2X^T X\theta - 2X^T y}{2m} \right) = 0$$

$$2X^T X\theta - 2X^T y = 0$$

$$\theta = \underline{\underline{(X^T X)^{-1} X^T y}}$$

b) Gradient descent is used to find optimal θ 's in a given linear equation of form

$$h(x) = \theta_0 x_0 + \dots + \theta_2 x_2 + \theta_1 x_1 + \theta_0 x_0$$

such that the cost function $J(\theta) = \frac{1}{2m} \sum_{i=1}^m h_0(x_i - y_i)^2$ is minimised

In order to minimise the above cost function, we need to calculate the slope of the function and adjust θ ~~in~~ such that we 'move down' the cost slope.

This is done by differentiating the cost function, which gives us

$$\frac{\partial J(\theta_0, \theta_1)}{\partial \theta_1} = \frac{1}{m} \sum_{i=1}^m (h_0(x_i) - y_i) \cdot x_i^1$$

$$\frac{\partial J(\theta_0, \theta_1)}{\partial \theta_0} = \frac{1}{m} \sum_{i=1}^m (h_0 x_i - y_i)$$

The above slopes are multiplied by a set learning rate α and subtracted by θ_0 and θ_1 .

This gives us

$$\theta_0 = \theta_0 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i)$$

$$\theta_1 = \theta_1 - \frac{\alpha}{m} \sum_{i=1}^m (h_{\theta}(x^i) - y^i) \cdot x^i$$

Both θ_0 and θ_1 are updated simultaneously until convergence.

- (1) Stochastic gradient descent is, for the most part, the same as batch gradient descent.

The major difference is in the way stochastic processes its iterations. Whereas batch gradient descent processes all datapoints at once, ~~stochastic~~ for each iteration, stochastic processes one row per iteration.

Formally its specified as follows.

~~while~~ while $|\text{new } \theta - \text{old } \theta| > \text{convergence}$:

{

for $i=0 : i < n ; i++$.

{ $\theta_{j=0 \text{ to } m} = \theta_j - \alpha (h_{\theta}(x^i) - y^i) x_j$

}

}