

# Problem 6

In [626]:

```
hw3 = pd.read_csv("hw3_dataset.txt", delim_whitespace=True, names = ['id', 'fund  
ing', 'fv', 'shares', 'l_buyout'])
```

## 6 A

Trying out EDA

In [630]:

```
hw3.describe()
```

Out[630]:

	funding	fv	shares	l_buyout
count	482.000000	4.820000e+02	4.820000e+02	482.000000
mean	0.439834	2.651672e+07	2.227942e+06	0.093361
std	0.496883	2.632174e+07	1.413872e+06	0.291240
min	0.000000	1.200000e+06	3.000000e+05	0.000000
25%	0.000000	1.025000e+07	1.300000e+06	0.000000
50%	0.000000	1.950000e+07	2.000000e+06	0.000000
75%	1.000000	3.250000e+07	2.700000e+06	0.000000
max	1.000000	2.346000e+08	1.101862e+07	1.000000

The fv and shares features have really large values compared to l\_buyout and funding columns.

In [627]:

```
hw3.head()
```

Out[627]:

	id	funding	fv	shares	l_buyout
0	1	0	1200000	3000000	0
1	2	0	1454000	1454000	1
2	3	0	1500000	300000	0
3	4	0	1530000	510000	0
4	5	0	2000000	800000	0

In [628]:

```
"""  
we know that id is not giving any useful information  
also, we'll check for missing values here  
"""
```

```
hw3 = hw3.drop('id', axis=1)  
nas = hw3.isnull().sum()  
print(nas)
```

```
funding      0  
fv           0  
shares       0  
l_buyout     0  
dtype: int64
```

There are no missing values in this dataset

In [632]:

```
features = hw3.drop('funding', axis=1)  
target = hw3['funding']
```

In [633]:

```
X_train, X_test, y_train, y_test = train_test_split(features, target, test_size=  
0.33, random_state=42)
```

In [634]:

```
X_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 322 entries, 157 to 102
Data columns (total 3 columns):
fv          322 non-null int64
shares      322 non-null int64
l_buyout    322 non-null int64
dtypes: int64(3)
memory usage: 10.1 KB
```

In [635]:


```
X_train.head()
```

Out[635]:

	fv	shares	l_buyout
<b>157</b>	13000000	1300000	0
<b>449</b>	63000000	4500000	1
<b>118</b>	10103125	1325000	0
<b>114</b>	9625000	1375000	0
<b>439</b>	55000000	2500000	0

In [653]:


```
sns.pairplot(hw3, vars=hw3.columns)
plt.show()
```



In [644]:

```
fig, axs = plt.subplots(1,3, figsize=(15, 5), facecolor='w', edgecolor='k')
fig.subplots_adjust(hspace = .5, wspace=.5)

for i,j in enumerate(X_train.columns):
    axs[i].hist(X_train[j])
    axs[i].set_title(j)
```




In [649]:

```
corr = hw3.corr()
mask = np.zeros_like(corr, dtype = np.bool)
mask[np.triu_indices_from(mask)] = True
plt.subplots(figsize = (10, 10))
sns.heatmap(corr, mask = mask, xticklabels = 1,
            yticklabels = 1, linewidths = 1, cbar_kws = {"shrink": .5})
```

Out[649]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f01443329b0>



**Inference**

In [654]:

```
X_train.head()
```

Out[654]:

	fv	shares	l_buyout
157	13000000	1300000	0
449	63000000	4500000	1
118	10103125	1325000	0
114	9625000	1375000	0
439	55000000	2500000	0

## 6 B

In [ ]:

```
X_train['capital'] = np.log(X_train['face_value'] * X_train['n_shares'])
X_validate['capital'] = np.log(X_validate['face_value'] * X_validate['n_shares'])
X_test['capital'] = np.log(X_test['face_value'] * X_test['n_shares'])

X_train['fv1'] = np.log(X_train.face_value)
X_train['nsl'] = np.log(X_train.n_shares)

X_validate['fv1'] = np.log(X_validate.face_value)
X_validate['nsl'] = np.log(X_validate.n_shares)

X_test['fv1'] = np.log(X_test.face_value)
X_test['nsl'] = np.log(X_test.n_shares)
```

The first thing to try is the log value of the two features which have a positively skewed distr.

In [660]:

```
X_train['fv_log'] = np.log(X_train['fv'])
X_test['fv_log'] = np.log(X_test['fv'])
X_train['shares_log'] = np.log(X_train['shares'])
X_test['shares_log'] = np.log(X_test['shares'])
```

/opt/conda/lib/python3.6/site-packages/ipykernel\_launcher.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

after removing the cwd from sys.path.

/opt/conda/lib/python3.6/site-packages/ipykernel\_launcher.py:5: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

"""

In [710]:

*# method to print basic metrics of a classifier*

```
def print_metrics(y_test, y_hat):
    print("Accuracy of model is ", metrics.accuracy_score(y_test, y_hat))
    print("Precision of model is ", metrics.precision_score(y_test, y_hat))
    print("Recall of model is ", metrics.recall_score(y_test, y_hat))
    print("Confusion matrix\n", metrics.confusion_matrix(y_test, y_hat))
```

*# method to plot AUC*

```
def plot_roc_auc(model, X_test):
    y_hat_probabilities = model.predict_proba(X_test)[::,1]
    fpr, tpr, _ = metrics.roc_curve(y_test, y_hat_probabilities)
    auc = metrics.roc_auc_score(y_test, y_hat_probabilities)
    plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
    plt.legend(loc=4)
    plt.show()
```



In [711]:


```
lr = LogisticRegression()  
lr.fit(X_train[['fv_log', 'shares_log', 'l_buyout']], y_train)  
  
y_hat = lr.predict(X_test[['fv_log', 'shares_log', 'l_buyout']])  
print_metrics(y_test, y_hat)
```

```
Accuracy of model is 0.5625  
Precision of model is 0.43333333333333335  
Recall of model is 0.41935483870967744  
Confusion matrix  
[[64 34]  
 [36 26]]
```

So, that's an okayish accuracy. We just tried all variables we thought would work. Let's try plotting the Area under curve.

In [712]:

```
plot_roc_auc(lr, X_test[['fv_log', 'shares_log', 'l_buyout']])
```



Let's try more feature transformation. Since face\_value and number of shares are highly correlated in stock market, it's worth trying. Moreover, there's no point in using two highly correlated features in building a model. the correlation between them is greater than 0.8. I will multiply them and then take a log over the multiplied quantity.

In [713]:

```
# All the values are greater than 0 so no need to do log(p+1) as shown in class.
```

```
X_train['stock_value'] = np.log((X_train['fv'] * X_train['shares']))
X_test['stock_value'] = np.log((X_test['fv'] * X_test['shares']))
```

/opt/conda/lib/python3.6/site-packages/ipykernel\_launcher.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>  
after removing the cwd from sys.path.

In [714]:

```
lr = LogisticRegression()
lr.fit(X_train[['stock_value', 'l_buyout']], y_train)

y_hat = lr.predict(X_test[['stock_value', 'l_buyout']])
print_metrics(y_test, y_hat)
```

Accuracy of model is 0.55625

Precision of model is 0.0


Recall of model is 0.0

Confusion matrix

```
[[89  9]
 [62  0]]
```

In [715]:

```
plot_roc_auc(lr, X_test[['stock_value', 'l_buyout']])
```



Didn't give a significant improvement. Logistic regression has different solver so will try a different solver.

In [736]:

```
lr = LogisticRegression(solver='lbfgs')
lr.fit(X_train[['stock_value', 'l_buyout']], y_train)

y_hat = lr.predict(X_test[['stock_value', 'l_buyout']])
print_metrics(y_test, y_hat)
```

Accuracy of model is 0.55

Precision of model is 0.4074074074074074


Recall of model is 0.3548387096774194

Confusion matrix

```
[[66 32]
 [40 22]]
```

In [717]:

```
plot_roc_auc(lr, X_test[['stock_value', 'l_buyout']])
```



Though it didn't improve significantly, it is definitely better than the previous model.

## 6 C

Now, we try another kind of classifier- LDA. Again following a similar approach, let's start with all features and slowly narrow down to the ones which we found useful from our logistic regression model.

In [718]:

```
lda = sklearn_lda()  
lda.fit(X_train[['fv_log', 'shares_log', 'l_buyout']], y_train)  
  
y_hat = lda.predict(X_test[['fv_log', 'shares_log', 'l_buyout']])  
print_metrics(y_test, y_hat)
```

```
Accuracy of model is  0.575  
Precision of model is  0.45  
Recall of model is  0.43548387096774194  
Confusion matrix  
[[65 33]  
 [35 27]]
```

Now, trying the enhanced version:


In [719]:

```
lda = sklearn_lda()  
lda.fit(X_train[['stock_value', 'l_buyout']], y_train)  
  
y_hat = lda.predict(X_test[['stock_value', 'l_buyout']])  
print_metrics(y_test, y_hat)
```

```
Accuracy of model is  0.55  
Precision of model is  0.4074074074074074  
Recall of model is  0.3548387096774194  
Confusion matrix  
[[66 32]  
 [40 22]]
```

In [720]:

```
plot_roc_auc(lda, X_test[['stock_value', 'l_buyout']])
```



In [731]:


```
lda = sklearn_lda(solver='svd')
lda.fit(X_train[['stock_value', 'l_buyout']], y_train)

y_hat = lda.predict(X_test[['stock_value', 'l_buyout']])
print_metrics(y_test, y_hat)
```

```
Accuracy of model is 0.55
Precision of model is 0.4074074074074074
Recall of model is 0.3548387096774194
Confusion matrix
[[66 32]
 [40 22]]
```

In [732]:

```
plot_roc_auc(lda, X_test[['stock_value', 'l_buyout']])
```



Let's even test the training accuracy of our model

In [723]:

```
y_hat_train = lda.predict(X_train[['stock_value', 'l_buyout']])
print_metrics(y_train, y_hat_train)
```

```
Accuracy of model is 0.5807453416149069
Precision of model is 0.5619834710743802
Recall of model is 0.4533333333333333
Confusion matrix
[[119 53]
 [ 82 68]]
```

Here, we observe again that feature transformation provides better AUC compared to fitting the model on the original features. However, changing solver here didn't make any difference.

## 6 D

Let's try finding training accuracies:

a) Logistic regression


In [725]:

```
lr = LogisticRegression(solver='lbfgs')
lr.fit(X_train[['stock_value', 'l_buyout']], y_train)

y_hat_train = lr.predict(X_train[['stock_value', 'l_buyout']])
print_metrics(y_train, y_hat_train)

y_hat_probabilities = lr.predict_proba(X_train[['stock_value', 'l_buyout']])[:,
1]
fpr, tpr, _ = metrics.roc_curve(y_train, y_hat_probabilities)
auc = metrics.roc_auc_score(y_train, y_hat_probabilities)
plt.plot(fpr, tpr, label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

Accuracy of model is 0.5745341614906833  
Precision of model is 0.5546218487394958  
Recall of model is 0.44  
Confusion matrix  
[[119 53]  
 [ 84 66]]




In [727]:

```
lr = LogisticRegression(solver='lbfgs')
lr.fit(X_train[['fv_log', 'l_buyout', 'shares_log']], y_train)

y_hat_train = lr.predict(X_train[['fv_log', 'l_buyout', 'shares_log']])
print_metrics(y_train, y_hat_train)

y_hat_probabilities = lr.predict_proba(X_train[['fv_log', 'l_buyout', 'shares_log']])[:,1]
fpr, tpr, _ = metrics.roc_curve(y_train, y_hat_probabilities)
auc = metrics.roc_auc_score(y_train, y_hat_probabilities)
plt.plot(fpr, tpr, label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

Accuracy of model is 0.5838509316770186  
Precision of model is 0.5634920634920635  
Recall of model is 0.47333333333333333  
Confusion matrix  
[[117 55]  
 [ 79 71]]



b) Linear Discriminant Analysis

In [737]:


```
lda = sklearn_lda(solver='svd')
lda.fit(X_train[['stock_value', 'l_buyout']], y_train)

y_hat_train = lda.predict(X_train[['stock_value', 'l_buyout']])
print_metrics(y_train, y_hat_train)

y_hat_probabilities = lda.predict_proba(X_train[['stock_value', 'l_buyout']])[:,1]
fpr, tpr, _ = metrics.roc_curve(y_train, y_hat_probabilities)
auc = metrics.roc_auc_score(y_train, y_hat_probabilities)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

Accuracy of model is 0.5807453416149069  
Precision of model is 0.5619834710743802  
Recall of model is 0.45333333333333333  
Confusion matrix

```
[[119  53]
 [ 82  68]]
```






In [738]:

```
lda = sklearn_lda(solver='svd')
lda.fit(X_train[['fv_log', 'l_buyout', 'shares_log']], y_train)

y_hat_train = lda.predict(X_train[['fv_log', 'l_buyout', 'shares_log']])
print_metrics(y_train, y_hat_train)

y_hat_probabilities = lda.predict_proba(X_train[['fv_log', 'l_buyout', 'shares_log']])[:,1]
fpr, tpr, _ = metrics.roc_curve(y_train, y_hat_probabilities)
auc = metrics.roc_auc_score(y_train, y_hat_probabilities)
plt.plot(fpr,tpr,label="data 1, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```

Accuracy of model is 0.5900621118012422  
Precision of model is 0.5671641791044776  
Recall of model is 0.5066666666666667  
Confusion matrix  
[[114 58]  
 [ 74 76]]



Here I tested out training accuracies in my model.