

Reference: Deeplearning.ai by Andrew Ng

A: The value of S1 would be 2 since there are 2 feature and S3 would be 1 since the output is binary.

```
In [23]: from collections import deque
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
%matplotlib inline
```

```
In [2]: def get_cost_w_reg(A3, Y, parameters, lambd):

    m = Y.shape[1]
    W1 = parameters["W1"]
    W2 = parameters["W2"]

    cross_entropy_cost = np.squeeze(- 1 * np.sum(np.multiply(Y, np.log(1 + np.multiply((1 - Y)

    L2_regularization_cost = lambd * (np.sum(np.square(W1)) + np.sum(np.s

    return cross_entropy_cost + L2_regularization_cost
```

```
In [11]: def init(layer_dims):

    parameters = {}

    for l in range(1, len(layer_dims)):

        parameters['W' + str(l)] = np.random.randn(layer_dims[l], layer_d
        parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))

    return parameters

def forward_prop_activation(A_prev, W, b, activation):

    Z = np.dot(W, A_prev) + b
    linear_cache = (A_prev, W, b)
```

```

if activation == "relu":

    A = np.maximum(Z, 0)
    activation_cache = Z

elif activation == "sigmoid":

    A = (1 + np.exp(-Z))**(-1)
    activation_cache = Z

elif activation == "tanh":

    activation_cache = Z
    e_z = np.exp(Z)
    e_nz = np.exp(-Z)
    A = (e_z - e_nz)/(e_z + e_nz)

cache = (linear_cache, activation_cache)

return A, cache

def forward_prop(X, parameters, activations):

    caches = []
    A = X
    L = len(parameters) // 2

    for l in range(1, L):

        A_prev = A
        activation = activations.popleft()
        A, cache = forward_prop_activation(A_prev, parameters["W" + str(l)],
        caches.append(cache)

    activation = activations.popleft()
    AL, cache = forward_prop_activation(A, parameters["W" + str(L + 1)],
    caches.append(cache)

    return AL, caches

```

In [12]: **def** **backprop**(dZ, cache, **lambda**):

```

A_prev, W, b = cache
m = A_prev.shape[1]

dW = (1/m) * np.dot(dZ, A_prev.T) + (lambda * W) / m
db = (1/m) * np.sum(dZ, axis=1, keepdims=True)
dA_prev = np.dot(W.T, dZ)

```

```

    return (dA_prev, dW, db)

def backprop_activation(dA, cache, activation, lambd):

    linear_cache, activation_cache = cache

    if activation == "relu":
        Z = activation_cache
        dZ = np.array(dA, copy=True)
        dZ[Z <= 0] = 0

    elif activation == "sigmoid":

        Z = activation_cache
        s = np.power((1 + np.exp(-Z)), -1)
        dZ = dA * s * (1 - s)

    elif activation == "tanh":
        Z = activation_cache
        s = (np.exp(Z) - np.exp(-Z)) / (np.exp(Z) + np.exp(-Z))
        A = Z * (1 - np.power(s, 2))

    dA_prev, dW, db = backprop(dZ, linear_cache, lambd)

    return (dA_prev, dW, db)

def get_backprop_gradients(AL, Y, caches, activations, lambd):

    grads = {}
    L = len(caches)
    m = AL.shape[1]
    Y = Y.reshape(AL.shape)

    dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))

    activation = activations.pop()
    current_cache = caches[L-1]

    grads["dA" + str(L-1)], \
    grads["dW" + str(L)], \
    grads["db" + str(L)] = backprop_activation(dAL, current_cache, activa

    for l in reversed(range(L-1)):

```

```

activation      = activations.pop()
current_cache = caches[l]

dA_prev_temp, dW_temp, db_temp = backprop_activation(grads["dA" + str(l) + "_prev"] +
                                                       current_cache)

grads["dA" + str(l)] = dA_prev_temp
grads["dW" + str(l + 1)] = dW_temp
grads["db" + str(l + 1)] = db_temp

return grads

```

```

def backward_w_reg(X, Y, cache, lambd):

    m = X.shape[1]
    (Z1, A1, W1, b1, Z2, A2, W2, b2) = cache

    dZ2 = A2 - Y
    dW2 = 1. / m * np.dot(dZ2, A1.T) + (lambd * W2) / m
    db2 = 1. / m * np.sum(dZ2, axis=1, keepdims=True)

    dA1 = np.dot(W2.T, dZ2)

    dZ1 = np.multiply(dA1, np.int64(A1 > 0))
    dW1 = 1. / m * np.dot(dZ1, X.T) + (lambd * W1) / m
    db1 = 1. / m * np.sum(dZ1, axis=1, keepdims=True)

    gradients = {"dA2": dA2, "dZ2": dZ2, "dW2": dW2, "db2": db2, "dA1": dA1,
                  "dZ1": dZ1, "dW1": dW1, "db1": db1}

    return gradients

```

```

In [13]: def update(parameters, grads, learning_rate):

    L = len(parameters) // 2

    for l in range(L):

        parameters["W" + str(l+1)] = parameters["W" + str(l+1)] - \
            learning_rate * grads["dW" + str(l+1)]
        parameters["b" + str(l+1)] = parameters["b" + str(l+1)] - \
            learning_rate * grads["db" + str(l+1)]

    return parameters

def train_model(X, Y, config):

    parameters = init(config["layers_dims"])

```

```

for i in range(0, config["num_iterations"]):

    fact          = deque(config["activations"].copy())
    AL, caches    = forward_prop(X, parameters, fact)
    cost          = get_cost_w_reg(AL, Y, parameters, config["lamdb"])
    bact         = deque(config["activations"].copy())
    grads         = get_backprop_gradients(AL, Y, caches, bact, config["
    parameters    = update(parameters, grads, config["learning_rate"])

    if i % 1000 == 0: print ("Cost after iteration %i: %f" % (i, cost)

return parameters

def predict(X, y, parameters, activations):

    p          = np.zeros((1,m))
    facts      = deque(activations.copy())
    probas, caches = forward_prop(X, parameters, facts)

    for i in range(0, probas.shape[1]):

        if probas[0,i] > 0.5: p[0,i] = 1
        else:                p[0,i] = 0

    return p

def load():

    input_data    = pd.read_csv('data/HW5.csv')
    train, test   = train_test_split(input_data, test_size = 0.2)

    x_train       = np.array(train.loc[:, ['a', 'b']])
    y_train       = np.array(train.loc[:, ['label']])

    x_test        = np.array(test.loc[:, ['a', 'b']])
    y_test        = np.array(test.loc[:, ['label']])

    return x_train, y_train, x_test, y_test

```

```
In [14]: x_train, y_train, x_test, y_test = load()
```

B. The accuracy for $S = 2$ is 73% with lambda 0.5 and 97% with lambda = 500

```
In [15]: config = { "activations": ["relu", "sigmoid"], "layers_dims" : [2, 2, 1],
                  "learning_rate": 0.0075, "num_iterations" : 3000,
                  "print_cost": True, "lambda": 0.5}

parameters      = train_model(x_train.T, y_train.T, config)
y_test_pred     = predict(x_test.T, y_test.T, parameters, config.get("a
accuracy_score(y_test, y_test_pred)
```

```
Cost after iteration 0: 0.693162
Cost after iteration 1000: 0.391115
Cost after iteration 2000: 0.359117
```

```
Out[15]: 0.747
```

```
In [16]: parameters
```

```
Out[16]: {'W1': array([[ -0.49614893, -0.77022891],
                      [-0.63621609, -0.32190759]]),
          'W2': array([[ -0.91698623,  0.7138217  ]]),
          'b1': array([[ 0.03989014],
                      [-0.03534747]]),
          'b2': array([[0.08445209]])}
```

lambda = 500

```
In [29]: config = {"activations": ["relu", "sigmoid"], "layers_dims" : [2, 2, 1],
                  "learning_rate": 0.0075, "num_iterations" : 3000,
                  "print_cost": True, "lambda": 500}

parameters      = train_model(x_train.T, y_train.T, config)
y_test_pred     = predict(x_test.T, y_test.T, parameters, config.get("a
accuracy_score(y_test, y_test_pred)
```

```
Cost after iteration 0: 0.693319
Cost after iteration 1000: 0.523287
Cost after iteration 2000: 0.464413
```

```
Out[29]: 0.979
```

In [28]: parameters

```
Out[28]: {'W1': array([[ -0.24860265, -0.5123004 ],
                    [ 0.75084355, -0.00124395]]),
          'W2': array([[ -0.5698433 ,  0.75442722]]),
          'b1': array([[0.02775391],
                    [0.08547043]]),
          'b2': array([[0.0220922]])}
```

C. The accuracy for $S = 10$ is 100% with $\lambda = 0.5$ and 74% with $\lambda = 500$

```
In [18]: config = {
    "activations": ["relu", "sigmoid"],
    "layers_dims" : [2, 10, 1], # 2-layer model
    "learning_rate": 0.0075,
    "num_iterations" : 3000,
    "print_cost": True,
    "lambda": 0.5
}

parameters = train_model(x_train.T, y_train.T, config)
y_test_pred = predict(x_test.T, y_test.T, parameters, config.get("activation"))
accuracy_score(y_test, y_test_pred)
```

```
Cost after iteration 0: 0.693806
Cost after iteration 1000: 0.049605
Cost after iteration 2000: 0.013739
```

```
Out[18]: 0.999
```

```
In [19]: parameters
```

```
Out[19]: {'W1': array([[ -0.55077971, -0.76099436],
                      [-0.36152371, -0.09346549],
                      [-0.17603503, -0.24320444],
                      [-0.28251728, -0.07297098],
                      [-0.08262698, -0.02141725],
                      [-0.65783904, -0.16983369],
                      [ 0.09517558,  0.24942912],
                      [ 0.20888765,  0.54431803],
                      [ 0.95332922,  0.15967837],
                      [ 0.17114534,  0.02866454]]),
          'W2': array([[ -0.94035605,  0.37418498, -0.30047745,  0.29224385,  0.
08449863,
                      0.68066009, -0.266931  , -0.58308542,  0.97117605,  0.174070
18]]),
          'b1': array([[ 0.04313984],
                      [-0.01973303],
                      [ 0.01346816],
                      [-0.01657169],
                      [-0.00331261],
                      [-0.0399877 ],
                      [-0.00238802],
                      [-0.00610014],
                      [ 0.09658692],
                      [ 0.01705597]]),
          'b2': array([[ -0.01104839]])}
```

Setting a high lambda


```
In [33]: config = {
    "activations": ["relu", "sigmoid"],
    "layers_dims" : [2, 10, 1], # 2-layer model
    "learning_rate": 0.0075,
    "num_iterations" : 3000,
    "print_cost": True,
    "lambda": 500
}

parameters = train_model(x_train.T, y_train.T, config)
y_test_pred = predict(x_test.T, y_test.T, parameters, config.get("activation"))
accuracy_score(y_test, y_test_pred)
```

```
Cost after iteration 0: 0.695011
Cost after iteration 1000: 0.654314
Cost after iteration 2000: 0.651400
```

```
Out[33]: 0.747
```

```
In [35]: parameters
```

```
Out[35]: {'W1': array([[ -1.95040110e-05,  8.07721180e-06],
    [-5.30872339e-02, -8.83093145e-02],
    [-7.65599550e-14,  3.18206948e-12],
    [-2.07695421e-03, -3.43965044e-03],
    [-2.38263393e-03, -3.95139184e-03],
    [-4.03199479e-06,  1.66886207e-06],
    [ 1.33921391e-06,  3.90191742e-07],
    [ 1.68743713e-03, -1.55409451e-04],
    [-1.22472288e-01, -2.03752742e-01],
    [-4.86712735e-05,  2.01483870e-05]]),
  'W2': array([[ 1.60461971e-05, -1.05405964e-01, -1.37457003e-12,
    -4.10613407e-03, -4.71700442e-03,  3.31758276e-06,
    1.92446883e-06,  1.98258249e-03, -2.43196958e-01,
    4.00460211e-05]]),
  'b1': array([[ -2.60617973e-04],
    [ 1.15538398e-01],
    [-2.35373103e-04],
    [ 4.29141017e-03],
    [ 5.01316378e-03],
    [-5.38563715e-05],
    [ 5.05891603e-04],
    [ 8.57900171e-03],
    [ 2.66827197e-01],
    [-6.50194286e-04]]),
  'b2': array([[0.31134879]])}
```

Part D

The data is not linearly separable. Part (ii) with 2 neurons in layer 2 could reach 74% accuracy where as part (iii) with 10 neurons in layer 3 could classify everything correctly. This makes sense as the more neurons there are, the higher the variance will be and more flexible the model will be.

```
In [26]: input_data = pd.read_csv('data/HW5.csv')
color    = ['blue' if i == 1 else 'red' for i in input_data['label']]

plt.scatter(input_data['a'], input_data['b'], color=color, alpha = 0.8)
```

Out[26]: <matplotlib.collections.PathCollection at 0x7fde5f1caf28>

